

Babelfish Compass: User Guide

Document history:

Jun-2023	v.2023-06: added automatic DDL generation; expanded scope of -optimistic option
Mar-2023	v.2023-03: added predefined 'optimistic' user-defined .cfg file and -optimistic option
Dec-2022	v.2022-12: added executive summary, analysis of dynamic SQL, reporting of complexity score for not-supported items, -userconfigfile option
Oct-2022	v.2022-10: added support for processing queries captured through extended events
Sep-2022	v.2022-09: generating .csv file for complexity/effort estimation
Jun-2022	v.2022-06-a: added -importfmt option to process captured query files; various report enhancements
Jun-2022	v.2022-06: added -pgimporttable option
Apr-2022	v.2022-04: added automatic check for new Compass version and -noudatechk option
Mar-2022	v.2022-03: added -recursive , -include , -exclude options
Feb-2022	v.2022-02: new Compass version numbering
Jan-2022	added compatibility matrix with Babelfish
Dec-2021	v.1.2: added -rewrite option
Nov-2021	correct typo in section about BabelfishCompassUser.cfg ; edit for grammar
Nov-2021	v.1.1: added user-definable overrides, example for -pgimport , Mac/Linux support
Oct-2021	v.1.0: first version

Contents

What Is Babelfish Compass?	3
Compatibility with Babelfish for PostgreSQL	4
Installing Babelfish Compass.....	5
Prerequisites.....	5
Downloading Babelfish Compass	5
Installation.....	5
Running Babelfish Compass on Windows.....	6
Running Babelfish Compass (Mac/Linux).....	7
Reports, applications, and input files.....	8
Report directory location	8
Specifying the Babelfish version	9
Command-line options.....	9
Examples	13
Automatic rewriting of unsupported features.....	15
File handling	16
The BabelfishFeatures.cfg file	17

SQL feature classifications.....	17
Example: BabelfishFeatures.cfg.....	17
The BabelfishCompassUser.cfg file (classification overrides)	19
Example: overriding default classification and reporting group	20
Predefined 'Optimistic' .cfg file	20
User-defined estimates in .csv file	22
Complexity estimates	22
Effort estimates	22
Examples	23
Using -pgimport	24
Schema for imported items.....	25
Example query.....	26
Processing captured SQL queries	27
SQL Server Profiler	27
SQL Server Extended Events.....	27
Examples	28
Automatic DDL generation	29
Command-line options.....	29
How it works.....	30
Generating only DDL	30
Example	32
Security.....	33
The -pgimport option	33
The -sqlpasswd option.....	33
Automatic update check	34
Using Babelfish Compass to migrate to PostgreSQL	35
Troubleshooting	37
Licensing.....	37

Document location

This document is located at https://raw.githubusercontent.com/babelfish-for-postgresql/babelfish_compass/main/BabelfishCompass_UserGuide.pdf .

What Is Babelfish Compass?

The Babelfish Compass tool (short for “**COMP**atibility **ASS**essment”) analyzes SQL/DDL code for one or more Microsoft SQL Server databases to identify the SQL features which are not compatible with Babelfish for PostgreSQL.

You can use Babelfish Compass to analyze the SQL/DDL code for your current SQL Server-based applications for compatibility with Babelfish. The purpose of this analysis is to gather information so you can make a Go/No Go decision about starting a migration project from SQL Server to Babelfish. For this purpose, Babelfish Compass produces an assessment report which lists -in great detail- all of the SQL features found in your SQL/DDL code, and whether or not these are supported by the latest version of Babelfish.

A new version of Babelfish Compass will be available as part of each Babelfish release containing new or changed functionality.

Note that Babelfish Compass is a stand-alone, on-premises tool. While Babelfish Compass is part of the Babelfish product, it is technically separate from Babelfish itself as well as from the Babelfish code, and is located in a separate GitHub repository.

Compatibility with Babelfish for PostgreSQL

The Babelfish Compass tool supports the following Babelfish versions.

In principle, any version of Babelfish Compass will support whichever Babelfish version the **BabelfishFeatures.cfg** file has been updated for. However, a full version of Babelfish Compass, also including fixes and enhancements, will in principle be published for every Babelfish release with added support for new T-SQL features.

Babelfish Compass version	Supports new Babelfish versions (also supports previous versions mentioned below)
2023-06	3.2.x (PG 15) 2.5.x (PG 14) 1.6.x, 1.5.x (PG 13)
2023-03[-a]	3.1.x (PG 15) 2.4.x (PG 14)
2022-12	2.3.x (PG 14) 1.4.x (PG 13)
2022-09, 2022-10, 2022-11	2.2.x (PG 14)
2022-07 2022-06[-a]	2.1.x (PG 14) 1.3.x (PG 13)
2022-04	same as below
2022-03	1.2.x (PG 13)
2022-02	same as below
1.2	1.1.x (PG 13)
1.0, 1.1	1.0.x (PG 13)

In February 2022, Babelfish Compass changed to a different version numbering schema (YYYY-MM) to avoid confusion with Babelfish version numbers. Consequently, Compass version 1.2 was followed by version 2022-02.

Note that there are no Babelfish version 2.0.x and 3.0.x.

Installing Babelfish Compass

Prerequisites

Before installing Babelfish Compass, you must install a Java Runtime Environment (JRE) version 8 or higher (64-bit version).

Babelfish Compass produces compatibility assessment reports in HTML format. To view the HTML output, we recommend using a recent release of the Google Chrome or Mozilla Firefox browser.

On Mac/Linux, you need to be able to run a **bash** script (e.g. with **#!/bin/bash**).

Downloading Babelfish Compass

Babelfish Compass is available as an open-source project at https://github.com/babelfish-for-postgresql/babelfish_compass.

A binary version can be downloaded from:

https://github.com/babelfish-for-postgresql/babelfish_compass/releases/latest ; choose the most recent **BabelfishCompass_<version>.zip** file.

The installation instructions that follow are based on this version.

Installation

Babelfish Compass is distributed as an executable JAR file, which requires no CLASSPATH settings. The only environmental requirement is that the Java JRE is in the PATH.

Installation steps on Windows:

1. Download the **BabelfishCompass.zip** file as detailed in the previous section.
2. Unzip the file so that the contents are placed in your installation directory of choice; this document will assume the file resides in **C:\BabelfishCompass**.
 - a. Do not install Babelfish Compass into **C:\Users\username\Documents\BabelfishCompass** since this is where the generated reports will be placed, and this location should be kept separate from the installation directory; Babelfish will not run when installed in this location.
 - b. If a previous installation is already present in your installation directory, you can overwrite the installation (but we recommend you make a backup copy first).
3. Installation is complete.

Installation steps on Mac/Linux:

1. Download the **BabelfishCompass.zip** file as detailed in the previous section.
2. Unzip this file so that the contents are placed in your directory of choice, for example **/home/username/BabelfishCompass** (Linux) or **/Users/username/BabelfishCompass** (Mac)
 - a. Do not install Babelfish Compass into **/home/username/BabelfishCompassReports** (Linux) or **/Users/username/BabelfishCompassReports** (Mac), since this is where the generated reports will be placed, and this location should be kept separate from the installation directory; Babelfish will not run when installed in this location.
 - b. If a previous installation is already present in your installation directory, you can overwrite the installation (but we recommend you make a backup copy first).
3. Verify the **BabelfishCompass.sh** shell script is executable by running **./BabelfishCompass.sh** . If it is not executable, run the command: **chmod +x BabelfishCompass.sh** .
4. Installation is complete.

Running Babelfish Compass on Windows

To run Babelfish Compass on Windows, open a **cmd** prompt (a "DOS box") and navigate to the Babelfish Compass installation directory.

Then, select the command line options that you need to include when invoking Babelfish Compass. The command line options are detailed in the Command-line options section of this guide, or you can review them on the command line by running:

```
C:\BabelfishCompass> BabelfishCompass[.bat] -help
```

Then, invoke **BabelfishCompass[.bat]** with your choice of command-line options.

Babelfish Compass usage typically starts by creating an assessment report file. The assessment report output file provides a detailed summary of the supported and unsupported SQL features in Babelfish for the analyzed SQL Server script(s). In the simplest usage case, a single SQL/DDI script is analyzed. To analyze a single script, simply specify a report name and an input file with your call to Babelfish Compass. For example:

```
C:\BabelfishCompass> BabelfishCompass[.bat] MyFirstReport C:\temp\AnyCompany.sql
```

This command creates an assessment report named **MyFirstReport**, containing the analysis for SQL/DDI script **AnyCompany.sql**.

When a report is created, BabelfishCompass will automatically:

1. Open an explorer window in the directory where the report files are stored.
2. Open the generated assessment report in the default browser.
3. Print the full pathname of the report file to **stdout**.

There are many additional command-line options you can include that support functionality to process multiple input scripts (for one application or multiple applications), generate more detailed output reports, and so on. See [Command-line options](#) for details.

Running Babelfish Compass (Mac/Linux)

To run Babelfish Compass on Linux, open a **bash** command prompt and navigate to the Babelfish Compass installation directory.

Then, select the command line options that you need to include when invoking Babelfish Compass. The command line options are detailed in the Command-line options section of this guide, or you can review them on the command line by running:

```
$ ./BabelfishCompass.sh -help
```

Then, invoke **BabelfishCompass.sh** with your choice of command-line options.

Babelfish Compass usage typically starts by creating an assessment report file. The assessment report output file provides a detailed summary of the supported and unsupported SQL features in Babelfish for the analyzed SQL Server script(s). In the simplest usage case, a single SQL/DDL script is analyzed. To analyze a single script, simply specify a report name and an input file with your call to Babelfish Compass. For example:

```
$ ./BabelfishCompass.sh MyFirstReport /tmp/AnyCompany.sql
```

This command creates an assessment report named **MyFirstReport**, containing the analysis for SQL/DDL script **AnyCompany.sql**.

When a report is created, BabelfishCompass will automatically:

1. Open a file browser in the directory where the report files are stored.
2. Open the generated assessment report in the default browser. Please note that on Linux, the browser will not open automatically; instead, simply open the file manually.
3. Print the full pathname of the report file to **stdout**.

There are many additional command-line options you can include that support functionality to process multiple input scripts (for one application or multiple applications), generate more detailed output reports, and so on. See [Command-line options](#) for details.

Reports, applications, and input files

The Babelfish Compass tool generates a report with a user-specified report name. The report is the result of analyzing one or more SQL/DDL scripts. In the simplest case, a single SQL/DDL script is analyzed. Babelfish Compass also supports combined analysis of multiple input scripts and multiple applications.

Each input script is associated with an application name. By default, the application name is taken from the input script file name. For example, a script named **Accounts.sql** is created for an application named **Accounts**. You can specify the application name with the **-appname** flag. A report can cover multiple input scripts for the same application, as well as multiple scripts for different applications.

Examples:

The following command generates a report for a single input file, with an application named **Accounts**:

```
BabelfishCompass MyReport C:\temp\Accounts.sql
```

The following command generates a report for a single input file, with an application named **Sales**:

```
BabelfishCompass MyReport C:\temp\ddl.20210913.sql -appname Sales
```

The following command generates a report for multiple input files, with an application named **Sales**:

```
BabelfishCompass MyReport C:\temp\ddl.20210913*.sql -appname Sales
```

The following command generates a report for multiple input files, with applications named **Accounts**, **Sales** and **HR**:

```
BabelfishCompass MyReport C:\temp\Accounts.sql C:\temp\Sales.sql C:\temp\HR.sql
```

When you create a report for multiple applications, the assessment can optionally indicate which applications contribute to a particular line item. To include this content, specify the **-reportoption apps** option. The report will contain lines in the following format:

```
SOUNDEX() : 45    #apps=3: Accounts(16), Support(20), HR(9)
```

This means 45 cases of the **SOUNDEX()** built-in function were found, in three applications as indicated:

Report directory location

By default, the Babelfish Compass report is created in the following location:

- **C:\Users\username\Documents\BabelfishCompass** (on Windows)
- **/Users/username/BabelfishCompassReports** (on Mac)

- **/home/username/BabelfishCompassReports** (on Linux)

A report is created as a **.html** file in a directory below this location. For example, on Windows, Babelfish Compass creates a report named **MyReport.html** in the following directory:

C:\Users\username\Documents\BabelfishCompass\MyReport.

Technically, the location of the report root directory is determined by the value of **System.getProperty("user.home")** in Java. The locations shown above are typical defaults. It is currently not possible specify a different, user-defined location instead.

Specifying the Babelfish version

By default, Babelfish Compass delivers a compatibility assessment for the most recent version of Babelfish, as indicated in the **BabelfishFeatures.cfg** file. You can perform an assessment for an earlier version of Babelfish by specifying the older version with the **-babelfish-version** option; for example:

-babelfish-version 2.2.0

The initial GA version of Babelfish was version 1.0.0. This is the oldest version you can specify. When the patch version (the third part in the version number) is omitted, **'0'** is assumed: specifying **2.2** is equivalent to **2.2.0**.

Command-line options

To display all of the command-line options, run **BabelfishCompass -help**. Note that all command-line options are optional:

- **-version**: displays the version of the Babelfish Compass tool.
- **-explain**: displays some high-level guidance on how to use the Babelfish Compass tool.
- **-encoding <encoding>**: specifies the encoding of the input files, if the files are not ASCII or the default encoding (this default is shown by **-help**).

The specified **-encoding** is applied to all input files. To process multiple input files with different encodings, import each file separately (with **-add**), specifying the correct encoding for each input file.

Unicode-formatted files with BOM bits are automatically detected and processed accordingly, so **-encoding** does not need to be specified.

To review a list of supported encodings, run **-encoding help**

- **-babelfish-version** <version>: performs the analysis for an older BBF version. See [Specifying the Babelfish version](#) for more information.
- **-add**: imports an additional SQL/DDL script to an existing report, performs an analysis, and generates a report.
- **-replace**: replaces an already-imported SQL/DDL script in an existing report, performs an analysis, and generates a report
- **-delete**: deletes all the files for an already existing report before recreating it.
- **-noreport**: performs an analysis without generating a report. This is useful when multiple files are imported; without **-noreport**, a report will be generated after every imported file. To generate a report after importing all files, include the **-reportonly** option.
- **-reportfile**: specifies the filename for the report. This does not affect the directory where the report files are located. See the [Examples](#).
- **-importonly**: imports the SQL/DDL script, but does not perform an analysis or generate a report. This can be useful when importing multiple files, as the analysis will otherwise be performed after every imported file.
- **-analyze**: performs an analysis on imported files, and generates a report. This can be used after importing files with **-importonly**, or to re-run an analysis on imported files from an earlier report (for example, when re-running the analysis when a later version of Babelfish has become available).
- **-userconfigfile** filename : as of v.2022-12, this specifies the user-defined **.cfg** file to be used (default = **BabelfishCompassUser.cfg**).
- **-optimistic** : as of v.2023-03, this is shorthand for specifying **-userconfigfile BabelfishCompassUser.Optimistic.cfg** (see [here](#)) plus **-rewrite**.
- **-sqlendpoint** , **-sqllogin**, **-sqlpasswd**, **-sqldblist** : see [Automatic DDL generation](#)
- **-nooverride**: do not use classification/report group overrides from the user-defined **.cfg** file.

- **-noreportcomplexity**: as of v.2022-12, estimated complexity for not-supported items will not be included in the Compass report.
- **-list**: displays the files/applications that have been imported for a report.
- **-reportonly**: generates a report for already-imported and analyzed SQL/DDI scripts. Specify a report name; do not specify input files. This option is useful when generating additional detailed assessment reports, for example with a cross-reference or additional filtering (see **-reportoption**).
- **-reportoption <options>**: specifies options for generating the final assessment report. Specify different options in a comma-separated list (without spaces), and/or by using multiple **-reportoption** flags. The cross-reference is not generated by default, as this potentially makes the assessment report very long.

Possible options are:

- **xref** or **xref=all**: generates two cross-references for all items that are marked as "not supported" or "review". One cross-reference is ordered by SQL feature, the other by objects for which such items were detected.

Warning: for large schemas, the report generated with **xref** (and even more so when combined with **status=all**), may become very large and may take longer to load in your browser. For this reason, the **xref** option is off by default, and you have to specify it explicitly with **-reportoption**.

In addition (v.2022-07 or later) object names are listed for which issues (or no issues) were identified by Compass. These lists can be reached by clicking on the "without issues" link in the Object Count section (without **xref**, this link is not present):

```
TABLE          : 112 (1012 columns)  (without issues: 109 of 112: list)
VIEW           :   1 (40 lines SQL)  (without issues: 1 of 1: list)
constraint FOREIGN KEY : 106
constraint PRIMARY KEY :  93
```

- **xref=feature** or **xref=object** generates only the cross-reference by feature, or by object, respectively.
- **status=<status>**: with **xref**, specifies the categories for which the cross-reference should be generated. Without this option, a cross-reference is generated only for items marked as "not supported" or "review". To generate a cross-reference for a different category, specify (for example) **status=supported** or **status=ignored**. With **status=all**, a

cross-reference for all items is generated.

Note that using this option can result in a longer assessment report.

- **detail**: with **xref**, generates additional detail for a reported item. For example, when reporting an object which cannot be created, specifying **detail** will include the name of the object. The report may get significantly longer as a result.
- **filter=<string>**: with **xref**, only includes items which match the specified string (case-insensitive). This can be useful when the generated cross-reference is very long, for example to cross-reference only specific items of interest. Note that the Summary section is not affected by this option.
- **linenrs=<number>**: with **xref**, defines the maximum number of line numbers mentioned in the cross-reference before suppressing the rest and adding "+ *NNN more*". By default, the maximum number is 10.
- **notabs**: with **xref**, opens the hyperlinks to the original SQL source code in the same browser window instead of in a new tab. By default, a hyperlink opens in a new tab. NB: For large SQL source files, it may take some time before the browser displays the desired line. If this takes too long, it is also possible to manually access the corresponding flat text file (same filename, but with a **.dat** suffix instead of **.html**).
- **apps**: shows which applications contribute to a particular line item in the Summary section when a report covers multiple applications. For example, the following means that 45 cases of the **SOUNDEX()** built-in function were found, in three applications as indicated:

SOUNDEX() : 45 #apps=3: Accounts(16), Support(20), HR(9)

- **batchnr**: with **xref**, displays the location of an item as a combination of the batch number in the file, the starting line number of the batch in the source file, and the line number in the batch. By default, the location is shown as the line number in the source file.
 - **hints**: (Compass v.2022-07 and later) lists all popup hints in the SQL Summary section at the end of the report (so that they can be read without requiring mouse action)
- **-quotedid {on|off}**: sets QUOTED_IDENTIFIER at the start of each SQL/DDl script. Default is ON
 - **-pgimport "pg-connection-attributes"**: creates a database table in a PostgreSQL database, and loads all captured items into the table. This table can then be accessed with SQL queries for further processing (see [Using -pgimport](#)). By default, this table is named **public.BBFCompass**, but a different name can be specified with **-pgimporttable**.

The PostgreSQL connection attributes are specified in a comma-separated list as follows:

host,port,username,password,dbname . The import is performed through a script created in the **captured** subdirectory. This script uses the PostgreSQL **psql** utility, which must be installed on your system, and in your PATH.

Note that the password is not saved anywhere and not written to any file (including temporary files).

- **-pgimportappend**: with **-pgimport**, appends content to an already-existing PostgreSQL table. Without this option, **-pgimport** will drop the table if it exists, before recreating it.
- **-pgimporttable**: with **-pgimport**, specifies the name of the table to import the data into.
- **-rewrite**: for certain T-SQL constructs that are currently unsupported by Babelfish, performs automatic rewriting of the applicable syntax with T-SQL features supported by Babelfish (see [Automatic rewriting of unsupported features](#)).
- **-exclude <list>** : specifies a comma-separated list of file type suffixes to be excluded. By default, a series of file types are excluded (unless overridden with **-include**); to display these, use **-help exclude**.
- **-include <list>** : specifies a comma-separated list of file type suffixes to be included; only the filetypes specified with **-include** will be processed.
- **-recursive** : any subsequent directory names are processed recursively, using all files in the directory tree as input. Both **-include** and **-exclude** (if specified) are applied to any files found. With **-recursive**, it is recommended to specify **-appname** as well, otherwise each input file will be assumed to represent a different application.
- **-noudatechk** : do not perform a check for a newer version of Babelfish Compass
- **-importfmt <format>** : process an XML file from SQL Server Profiler with captured SQL queries (see [Processing captured SQL queries](#))
- **-nodedup** : with **-importfmt**, do not perform de-duplication of captured SQL

Examples

Generate a default report without cross-references for an application named **Sales**:

```
BabelfishCompass MyReport C:\temp\Sales.sql
```

Generate a default report without cross-reference for an application named **Sales**, deleting the report directory first if it already exists:

```
BabelfishCompass MyReport C:\temp\Sales.sql -delete
```

Generate a report for applications named **Accounts** and **Sales**, cross-referencing all categories, including additional detail, and allowing up to 100 line numbers to be enumerated in the cross-reference:

```
BabelfishCompass MyReport2 C:\temp\account*.sql -appname Accounts -add -noreport
```

```
BabelfishCompass MyReport2 C:\temp\sales.sql -add -noreport
```

```
BabelfishCompass MyReport2 -reportoptions xref,status=all,detail,linenrs=100
```

Display all files and applications imported for MyReport2:

```
BabelfishCompass MyReport2 -list
```

Re-run an analysis for an existing report, but specifically for Babelfish version 1.5.0 (this example assumes the latest version of Babelfish is later than 1.5.0):

```
BabelfishCompass MyReport3 -analyze -babelfish-version 1.5.0
```

Import all captured items into a PostgreSQL database table:

```
BabelfishCompass MyReport3 -pgimport  
"mybigghost.anycompany.com,5432,bob,B!gbob72,mydb"
```

Generate a cross-referenced report named : **C:\...\BabelfishCompass\MyReport4\MyApp.xref.html**. (without the **-reportfile** option, the report file name would be something like

C:\...\BabelfishCompass\ MyReport4\report-MyReport4-2021-Sep-13-21.22.23.html):

```
BabelfishCompass MyReport4 C:\temp\MyApp.sql -reportfile MyApp.xref -reportoption xref
```

Generate a combined report for applications **Sales** and two applications **Finance** and **Inventory**, each of which consists of a directory tree containing **.sql** files on multiple levels, and perform automatic rewriting where possible:

```
BabelfishCompass MyReport5 -importonly C:\temp\Sales.sql
```

```
BabelfishCompass MyReport5 -add -importonly -appname Finance -recursive C:\Finance\install
```

```
BabelfishCompass MyReport5 -add -importonly -appname Inventory -recursive C:\Inventory\install
```

```
BabelfishCompass MyReport5 -analyze -rewrite -reportoption apps
```

Automatic rewriting of unsupported features

As of version 1.2 of Babelfish Compass, you can use the **-rewrite** option to address certain SQL features which are not currently supported by Babelfish, by rewriting the SQL feature in question in such a way that Babelfish is able to process it. One example is the MERGE statement.

- When not specifying the **-rewrite** option, the assessment report will include a section "**Automatic SQL Rewrite Opportunities**" which lists the SQL features that could be addressed with **-rewrite**, but without actually rewriting them.
- When specifying the **-rewrite** option, Babelfish Compass creates a subdirectory **rewritten** in the report directory, containing a copy of the original SQL source file in which specific features have been rewritten (if nothing is rewritten, no copy will be created in **rewritten**).
The assessment report will contain a section with the specific rewritten features.
When **-reportoption xref** is used, the cross-reference links in the 'rewritten' sections point to the rewritten SQL file (instead of to the original SQL file).

In a rewritten SQL file, the bottom of the file contains a list of all changes made by Babelfish Compass.

When using the **-rewrite** option, you should execute the rewritten SQL file against Babelfish instead of the original SQL file.

Notes:

- Using **-rewrite** may cause Babelfish Compass to run slower than without **-rewrite**, especially for large files in which many features are rewritten. For very large input files, it may therefore be practical to first run an analysis without **-rewrite**; when the Compass report indicates that rewrite opportunities were identified, then re-run Compass with the **-analyze -rewrite** flags.
- For dynamic SQL queries which contain unsupported-but-rewritable SQL features, no rewrite is performed. The SQL feature will be reported as 'Not Supported'.

File handling

An assessment report is an HTML file located in the report directory:

- On Windows: `%USERPROFILE%\BabelfishCompass\<report-name>`

A flat text version of the report is available in the same directory as the HTML file; this text version is named identically, but ends in **.txt** instead of **.html**.

The report directory contains multiple subdirectories as described below. You should not rename or edit the files in these subdirectories, as future invocations of Babelfish Compass for this report may no longer work correctly (or at all):

- **imported**: contains a copy of the original SQL/DDL input scripts. These are stored to allow re-running the analysis at a later time (for example, for a newer version of Babelfish). If the original input files used a specific encoding, the files in the imported directory are in UTF8 format.
For each imported file, an HTML version is also located in the **imported\html** directory. When generating a cross-reference in the assessment report, hyperlinks are generated to the actual line in the original document where the SQL feature was found.
- **imported\sym**: contains files with symbol table information, for Compass-internal use.
- **captured**: contains files that contain items that were captured during analysis. These are SQL features and options, which are reflected in the assessment report. The files in this directory can be imported into a PostgreSQL database table using the **-pgimport** option
- **log**: contains the session log file for each invocation of Babelfish Compass.
- **errorbatches**: is a directory created only when syntax errors were found in the imported SQL/DDL scripts. In this case, the input batches with the errors are saved in a file so that the user has access to this information. If desired, you can rename or delete these files as they are not used as input for any further processing steps.
- **rewritten**: contains rewritten input files as a result of using the **-rewrite** option. Only input files where actual rewriting was performed, will be present here.

The BabelfishFeatures.cfg file

The compatibility assessment performed by the Babelfish Compass tool is driven by the file **BabelfishFeatures.cfg**, which is located in the Babelfish Compass installation directory. This file contains definitions of features that are (not) supported in a specific Babelfish version.

For each Babelfish release containing changes in functionality, a new version of the **BabelfishFeatures.cfg** will also be released. When Babelfish Compass is already installed, the existing version of **BabelfishFeatures.cfg** should be replaced (overwritten) by the newer version of this file.

Do not edit, modify, or rename the **BabelfishFeatures.cfg** file; Babelfish Compass will detect changes, and terminate immediately.

SQL feature classifications

The general principle behind **BabelfishFeatures.cfg** is that features which are not listed in this file are supported by Babelfish. Features that are not supported may fall in either of these categories:

- **Not Supported** : the feature is currently not supported by Babelfish.
- **Review Semantics** : the feature involves aspects which cannot be addressed by Babelfish, but requires review to determine whether or not it requires changes to be made as part of the migration process.
- **Review Performance** : the feature involves a performance-related aspect in SQL Server, and therefore you should review this carefully to determine if performance may be impacted when running on Babelfish.
- **Review Manually** : the feature cannot be assessed by Babelfish Compass, but needs to be manually examined. For example: **SET LANGUAGE @v** : Babelfish Compass cannot determine if **@v** contains a Babelfish-supported language name.
- **Ignored** : the feature is currently ignored by Babelfish.

Example: BabelfishFeatures.cfg

The following example denotes that **ALTER VIEW** is not supported, and is reported in a group named **Views**:

```
[ALTER VIEW]
rule=create_or_alter_view
report_group=Views
```

The following example denotes that the only supported option for **FETCH** is **FETCH NEXT**; any **FETCH** options are reported in a group named **Cursors**:

```
[FETCH cursor]
rule=fetch_cursor
list=NEXT,PRIOR,FIRST,LAST,ABSOLUTE,RELATIVE
supported-1.0.0=NEXT
report_group=Cursors
```

For more information about the contents of **BabelfishFeatures.cfg**, see the file's header.

The BabelfishCompassUser.cfg file (classification overrides)

As described in the previous section, the **BabelfishFeatures.cfg** file defines which features are or are not supported in a particular version of Babelfish. For SQL features that are not supported, you can override the classification defined by **BabelfishFeatures.cfg**. For this purpose, Babelfish Compass generates a file named **BabelfishCompassUser.cfg**, which is located in the report root directory; see Report directory location for more information. The default location of this file is

C:\Users\username\Documents\BabelfishCompass\BabelfishCompassUser.cfg (on Windows). You can edit this file (unlike **BabelfishCompass.cfg**, which should not be modified by the user).

BabelfishCompassUser.cfg is not overwritten when installing a new version of Babelfish Compass, as opposed to **BabelfishFeatures.cfg**, which will always be replaced in a new version of Babelfish Compass.

In v.2022-12, the user-defined .cfg file to be used can be specified with **-userconfigfile**; the default remains **BabelfishCompassUser.cfg**.

Use of the **BabelfishCompassUser.cfg** file is not recommended for new users of Babelfish Compass. However, if you are an experienced user, you can use **BabelfishCompassUser.cfg** to tailor your assessment reports by putting more or less focus on specific SQL features.

The user-defined .cfg file allows you to:

- Override the classification of a not-supported SQL feature (e.g. 'Ignored' instead of 'Not Supported')
- Override the default complexity estimate (see next section)
- Define a user-defined effort estimate (see next section)

BabelfishCompassUser.cfg contains all of the sections that are present in **BabelfishFeatures.cfg**, like **[Datatypes]** or **[Built-in functions]**. You shouldn't modify these section headers, but can add certain items to a section, as described below.

Note that any modifications made to the **BabelfishCompassUser.cfg** will not be saved or stored by Babelfish Compass. Ensure that **BabelfishCompassUser.cfg** is properly backed up.

Babelfish Compass will create the **BabelfishCompassUser.cfg** file if it does not exist. If new sections have been defined in **BabelfishFeatures.cfg**, which are not yet in **BabelfishCompassUser.cfg**, the new sections will be appended. If you manually delete sections from **BabelfishCompassUser.cfg**, those section will be appended again the next time Babelfish Compass runs.

Note:

- User-defined overrides are applied during analysis, and any overridden values are recorded in the captured items; the assessment report is generated from these captured items. When only

generating a report (e.g. with **-reportonly**), no overrides will be applied.

When the user has modified override entries in **BabelfishCompassUser.cfg** and wants to apply this to a report, the **-analyze** flag should be used.

- When a user-defined override is applied, the captured items will reflect the values after the overrides have been applied; the original values have been lost. This means that it is not possible to determine for individual captured items whether an override was applied (for example, after using **-pgimport**).

Example: overriding default classification and reporting group

By default, **CLUSTERED** indexes and constraints are classified as **Review Semantics** by Babelfish Compass. If you decide you don't care about those aspects, and you want these to be ignored in the assessment report, the classification can be overridden in **BabelfishCompassUser.cfg** by adding **default_classification=Ignored** :

```
[CLUSTERED index]
default_classification=Ignored
```

Likewise, the **FORMAT()** and **STR()** functions are not supported in Babelfish version 1.0.0, and will be reported accordingly. If you want these functions to be classified as **Review Manually** and reported under **Formatting functions**, then add the following lines to the section **[Built-in functions]** in **BabelfishCompassUser.cfg** :

```
[Built-in functions]
default_classification=ReviewManually=FORMAT,STR
report_group=Formatting functions=FORMAT,STR
```

Note that these changes only affect how SQL features are classified in the Babelfish Compass report. There is no impact on how Babelfish itself processes the SQL features for which you changed the classification.

For more information about possible modifications that you can make to **BabelfishCompassUser.cfg**, see the file's header.

Predefined 'Optimistic' .cfg file

As of v.2023-03, a predefined file named **BabelfishCompassUser.Optimistic.cfg** is included with Babelfish Compass. After installing the new Babelfish Compass version, this file is copied from the installation directory to **C:\Users\username\Documents\BabelfishCompass** (on Windows) the next time Compass runs.

When specifying **-userconfigfile BabelfishCompassUser.Optimistic.cfg** , various T-SQL features which are unsupported by Babelfish will be reclassified as 'Ignored', thus removing them from the list of 'Not

Supported' features. This concerns T-SQL features that are unlikely to affect the actual application functionality and have a make-or-break effect on the success of migration the application, like ALTER DATABASE, or functions like FILEGROUP_NAME().

The reason for providing this .cfg file is that users are sometimes scared away from Babelfish by the initial Compass report which contains some of these 'false positive' T-SQL features. By removing these from the 'Not Supported' list, the Compass reports will present a more realistic -or, if you will, a more optimistic- view of how well the application qualifies for migration to Babelfish.

This predefined 'optimistic' .cfg file is intended primarily for such first impressions when no deep dive on the actual T-SQL feature in the application is performed. Once the actual migration is undertaken, there is no reason to use this .cfg file since all actual SQL aspects in the application will need to be handled regardless.

Compass users can modify **BabelfishCompassUser.Optimistic.cfg** as any other user-defined .cfg file (see above) but should be aware that each new version of Compass may include a new copy of this file, which will overwrite the existing copy. Therefore, if users want to customize

BabelfishCompassUser.Optimistic.cfg, they should rename this file and apply the customizations to the renamed copy – and use that copy with the **-userconfigfile** flag.

When using the **-optimistic** flag, this is equivalent to specifying **-userconfigfile BabelfishCompassUser.Optimistic.cfg** plus specifying **-rewrite**.

User-defined estimates in .csv file

Complexity estimates

As of Compass version 2022-09, Compass generates a **.csv** file with the same filename as the report file, to assist specialist Compass users in estimating the amount of work required to address not-supported items (as reported by Compass). The **.csv** file has a column 'Complexity' which, for each not-supported item, contains a value LOW, MEDIUM or HIGH reflecting a very rough complexity estimate for resolving the item in question.

The Compass user can override these default estimates by adding their own complexity scores in the **BabelfishCompassUser.cfg** file using the **complexity_score** keys (see examples below).

The **.csv** file is intended to be imported into a spreadsheet, and Compass user should add their own formulas to the spreadsheet for performing calculations.

It should be noted that the default Compass complexity estimates are generalized and somewhat arbitrary since the experience and expertise of the team performing the migration is a much more deciding factor for the overall effort required. These Compass-provided complexity estimates should therefore be taken as very rough high-level guidance only; Compass users are urged to evaluate and adjust all values in the context of the actual customer application being analyzed.

The user-specific override values in **BabelfishCompassUser.cfg** can be LOW, MEDIUM or HIGH, but also an arbitrary number between 0 and 100 (so as to allow Compass users to use more detail).

As of v.2022-12, the Compass report itself includes a summary of the complexity score for each item classified as 'Not Supported', displayed between square brackets:

```
BINARY_CHECKSUM() [medium] : 264
```

By specifying the **-noreportcomplexity** option, the complexity scores will not be included in the Compass report; the **.csv** file containing the complexity scores will however always be generated.

As of v.2022-12, when uploading analysis details into PG using the **-pgimport** option, the complexity score is available in column **misc**; user-defined effort estimates (see below) are not uploaded.

As of v.2023-06, user-defined effort estimates are available in column **misc2**.

Effort estimates

Apart from the Compass-provided complexity estimates, it is also possible to define your own effort estimates, expressed in minutes, hours or days, in the **BabelfishCompassUser.cfg** file using the **effort_estimate** keys (see examples below).

When user-defined effort estimates are used, these values will show up in two additional 'Effort' columns in the **.csv** file: the first column is a textual representation of the user-configured value (e.g.

'5 minutes' or '1 hour'), and the second is the corresponding number of minutes, for calculation purposes (e.g. 5 or 60).

Examples

The precise way of specifying these values is explained a bit more in the header of the **BabelfishCompassUser.cfg** file. The sections in this file correspond to the same section in **BabelfishFeatures.cfg** (which you cannot edit yourself), and in most cases correspond quite obviously to the grouping of reported items in the Compass report.

(NB. the actual values shown below are chosen arbitrarily and should not be used as guidance)

[Built-in functions]

```
complexity_score-HIGH=SOUNDEX    # complexity = HIGH for this function, if unsupported
complexity_score=LOW              # complexity = LOW for any other unsupported function
complexity_score-60=COL_LENGTH    # complexity = 60 for this function, if unsupported
```

[Cursor options]

```
effort_estimate-4hours=SCROLL,FOR_UPDATE    # 4 hours for these options, if unsupported
effort_estimate=1hour                       # 1 hour for any other unsupported options
```

Using -pgimport

As described earlier, the **-pgimport** flag lets you load all captured items into a PostgreSQL table. From here, you can perform customized additional operations on this data. Before you can use **-pgimport**, the [PostgreSQL psql client](#) needs to be installed on your system, and needs to be in the PATH.

By default, data is imported into a table named **public.BBFCompass**, but a different name can be specified with the **-pgimporttable** option.

Note: when using the **-pgimport** flag, any user-defined complexity estimates or effort estimates in the user-defined .cfg file will only be included in the uploaded data when the user-defined .cfg file is specified together with **-pgimport** (for the "optimistic" .cfg file, you can specify **-optimistic** instead).

Examples of what you may be able to do with **-pgimport**:

- Run SQL queries to find objects with a complex combination of attributes. For example: find all SQL functions with at least two parameters, including a MONEY-type parameter, a SMALLDATETIME result type, and a table variable operation in the function body.
- The Babelfish Compass assessment report deliberately does not report any 'compatibility percentage', because it is difficult to define such a number in a meaningful way. A simple way to calculate such a percentage would be to take the ratio of non-supported features vs. supported features. However, some unsupported features may be very difficult to work around while other may be easy. Yet, they would both weigh equally heavy in such a calculation.

You can decide how to calculate a viable compatibility percentage for your evaluation. For example, you could write a SQL-based application that assigns different weights to different non-supported features, thus calculating a more realistic compatibility percentage on the basis of the captured items that were loaded with **-pgimport**.

Note: any such calculations are the exclusive responsibility of the Babelfish Compass user.

- When a migration opportunity is discussed, a key question is to estimate the time and cost of performing a migration. While this question is realistic, the Babelfish Compass tool does not attempt to make any estimates with respect to the amount of time or effort it may require to address the non-supported issues that were identified. The reason is that the actual effort required will be highly dependent on skills and experience of the individuals doing the actual work (picture a team of seasoned DBAs with decades of database experience vs. a team of newly arrived university graduates). Since it is not realistic to generalize such effort estimates, Babelfish Compass does not attempt this.

However, you could try to build such functionality yourself on the basis of the captured items that were loaded with **-pgimport**. Imagine an experienced team of migration experts who have collected detailed data points from their past migration projects; such a team might be able to quantify the effort required for the non-supported items in the Babelfish Compass assessment

report, specifically aimed at their own team with their specific experience. The **-pgimport** function makes it possible to build an application using the imported items for making effort estimates.

Note: any such estimates are the exclusive responsibility of the Babelfish Compass user.

Schema for imported items

When you include the **-pgimport** flag, Babelfish Compass creates a PostgreSQL table with the following definition:

```
CREATE TABLE BBFCompass(  
    babelfish_version VARCHAR(20) NOT NULL,  
    date_imported TIMESTAMP NOT NULL,  
    item VARCHAR(200) NOT NULL,  
    itemDetail VARCHAR(200) NOT NULL,  
    reportGroup VARCHAR(50) NOT NULL,  
    status VARCHAR(20) NOT NULL,  
    lineNr INT NOT NULL,  
    appName VARCHAR(100) NOT NULL,  
    srcFile VARCHAR(300) NOT NULL,  
    batchNrInFile INT NOT NULL,  
    batchLineInFile INT NOT NULL,  
    context VARCHAR(200) NOT NULL,  
    subcontext VARCHAR(200) NOT NULL,  
    misc VARCHAR(20) NOT NULL,  
    misc2 BIGINT NOT NULL -- as of v.2203-06  
);
```

The columns represent the following:

- **babelfish_version** : is the Babelfish version for which analysis was performed.
- **date_imported** : is the date/time that **-pgimport** ran.
- **item** : is a line item as shown in the report.
- **itemDetail** : is additional info about a line item.
- **reportGroup** : is the report group as show in the report.
- **status** : is the classification of the item; for example, **SUPPORTED** or **NOTSUPPORTED**.
- **lineNr** : is the line number of the item in the T-SQL batch.
- **appName** : is the application name.
- **srcFile** : is the SQL source file name.
- **batchNrInFile** : is the batch number of the T-SQL batch in SQL source file.
- **batchLineInFile** : is the line number in the file at the start of the batch.

- context : is the name of an object, or T-SQL batch.
- subcontext : is the (optional) name of a table in the object.
- misc : as of v.2022-12: complexity score; in previous versions: not used.
- misc2 : added in v.2023-06: effort estimate in minutes, if specified.

Example query

You can run SQL queries against the imported items to locate specific information. For example, to find this information:

"find all SQL functions with at least two parameters, including a MONEY-type parameter, a SMALLDATETIME result type, and a table variable operation in the function body"

...use this SQL query:

```
select distinct context from BBFCompass
-- filter on table variable operation:
where item like '% @tableVariable'
-- filter on function with >= 2 parameters:
and context in (
    select context from BBFCompass
    where context like 'FUNCTION %'
    and item like '% parameter'
    group by context
    having count(*) >= 2)
-- filter on MONEY-type parameter:
and context in (
    select context from BBFCompass
    where context like 'FUNCTION %'
    and item like 'MONEY %function parameter%')
-- filter on function result type:
and context in (
    select context from BBFCompass
    where context like 'FUNCTION %'
    and item like 'SMALLDATETIME %scalar function result%')
```

Note: On large tables, performance may benefit from adding indexes to one or more columns of this table. This is left for the user to explore.

Processing captured SQL queries

Apart from server-side DDL, also client-side SQL queries should be considered during a database migration. When capturing client-side SQL queries as described below, Babelfish Compass can extract the SQL queries from the capture files and perform a Compass assessment on them.

In order to process capture files, specify the command-line option **-importfmt *format***, as shown below.

Since captured SQL often contains many near-duplicate statements that only differ in the value of a lookup key or a constant, by default Compass de-duplicates the captured SQL prior to analysis. De-duplication is performed by masking the values of all string/numeric/hex constants.

To suppress de-duplication, specify the command-line option **-nodedup**.

The extracted and de-duplicated SQL queries/batches are saved into a file in directory **extractedSQL**: a file named **MyCapture.xml** will be saved into **extractedSQL/MyCapture.xml.extracted.sql**. This file is then used as input for the Compass analysis.

SQL Server Profiler

To capture SQL statements with SQL Server Profiler, take these steps:

1. In SQL Server Profiler, under "Trace Properties", use the **TSQL_Replay** template
2. Initiate the tracing in SQL Server Profiler
3. Run the client application against the SQL Server database
4. When done capturing the client application's SQL, save the captured results (in SQL Server Profiler) with **Save As → Trace XML File for Replay**. This creates an XML file containing the captured SQL batches.
5. Run Babelfish Compass with the just-created XML file as input, and specify the command-line option **-importfmt MSSQLProfilerXML**
 - The extracted SQL batches are saved into a file in directory **extractedSQL**: a file named **MyCapture.xml** will be saved into **extractedSQL/MyCapture.xml.extracted.sql**

SQL Server Extended Events

To capture SQL statements with SQL Server Extended Events, take these steps:

1. Run the client application against the SQL Server database
2. Use SQL Server Extended Events to capture SQL queries
3. Extract the captured events from the **.xel** file as **.xml** files containing `<event...> ... </event>` XML documents. Note that the **.xel** files cannot be processed by Compass.
4. Run Babelfish Compass with the XML file as input, and specify the command-line option **-importfmt extendedEventsXML**

Examples

BabelfishCompass MyReport C:\temp\MyProfilerCapture.xml -importfmt MSSQLProfilerXML

BabelfishCompass MyReport C:\temp\MyXECapture.xml -importfmt extendedEventsXML

Automatic DDL generation

As of Babelfish Compass v.2023-06, Compass may optionally perform the DDL generation directly for the SQL Server database(s) for which analysis needs to be performed. This means that the Compass user does not have to manually generate the DDL script using SQL Server Management Studio. This may not always be the most optimal solution (see below), but it can simplify the overall process.

Command-line options

To let Babelfish Compass perform the DDL generation, the following command-line options must be specified. Compass will then connect to the SQL Server, generate the DDL file(s), and run a Compass analysis on the generated files:

- **-sqlendpoint** : the hostname or IP address of the SQL Server; optionally, the port number can also be specified (e.g. **10.123.45.67,1433** or **mybigbox,1433**).
- **-sqllogin** : the login name to connect to the SQL Server; this login should typically be a member of the **sysadmin** role so as to have permission to generate the DDL.
- **-sqlpasswd** : the corresponding password.
- **-sqldblist** : optional; when omitted or when 'all' is specified, DDL is generated for all user databases in the server; alternatively, a comma-separated list of database names can be specified.

Notes:

- While generating DDL, Compass creates a network connection to the SQL Server specified, through a Powershell script (see below).
- When one of options **-sqlendpoint** , **-sqllogin** or **-sqlpasswd** is specified, the others must be specified as well.
- These options cannot be combined with specifying input files; also, this cannot be used to add DDL files to an existing report with **-add** or **-replace**, so you must either create a new report or use the **-delete** flag. Other Compass options such as **-rewrite** , **-optimistic** and **-reportoption** can be specified as usual.
- One DDL file is generated per T-SQL user database. When more than one database is processed, the flag **-reportoption apps** is automatically applied so that the Compass report shows how the T-SQL features are distributed over the different databases.
- The DDL files are generated into a directory named similar to **CompassAutoDDL-2023-Jun-07-13.22.51** under your session's **%TEMP%** location (on Windows) or **/tmp** (Mac/Linux). The directory is reported by Compass on the console output. The directory is not deleted afterwards so the original DDL files are retained.
- The speed of generating the DDL is highly dependent on the network proximity to the SQL Server since this involves many client-server roundtrips. When a few thousand objects exist in

the database, it may take minutes to hours to complete, depending on the network proximity. You can tell that DDL generation has completed when the message '**DDL generated in <directory>**' occurs, and Compass starts its analysis of the generated DDL.

If it takes an unreasonably long time to generate the DDL, then you might want to consider to use SQL Server Management Studio instead.

- Even with optimal network connectivity, generating the DDL will likely take more time than analyzing it by Compass.
- On Windows, if the password contains a ^ character (a.k.a. 'caret', 'circumflex', or SHIFT-6), enclose the password in double quotes since this is an escape character in Windows .bat and Powershell.

How it works

Babelfish Compass uses .Net SQL Management Objects (SMO) to perform the DDL generation. Compass invokes a Powershell script named **SMO_DDL.ps1**, which is located in the Babelfish installation directory. This script uses SMO and connects to the SQL Server specified by the user, requiring a network connection.

On Windows, Powershell is available by default. On Linux and Mac, Powershell first needs to be installed by the user before running Compass (see <https://learn.microsoft.com/en-us/powershell/scripting/install/installing-powershell-on-linux>) .

When Powershell script **SMO_DDL.ps1** is invoked, it will install SMO if needed (typically this only happens when it runs for the first time); this installation may take some time, and requires a network connection.

This Powershell script, and the installation of SMO, have been tested successfully on Windows, Linux and Mac. However, in rare cases installing SMO might fail. The remedy would be to install SMO manually but that may not be trivial given that the installation steps in the script have already failed (this user guide does not provide guidance how to install SMO manually). Alternatively, use the manual approach to generate DDL via SQL Server Management Studio.

Generating only DDL

The script **SMO_DDL.ps1** can also be used stand-alone to generate DDL. This can be useful when there are many SQL Server instances and/or many databases per instance, and perhaps you want to inspect or consolidate the DDL scripts before running a Compass analysis on them. For instructions how to run the script stand-alone, see the header of the script.

Notes:

- Do not make any changes to the **SMO_DDL.ps1** script, since Compass expects this script to be available without any changes. If you want to make changes, take a copy and modify the copy.
- As above, Powershell needs to be installed and available.

Example

```
C:\BabelfishCompass> BabelfishCompass.bat SMOTest -sqlendpoint mybigbox -sqllogin
sa -sqlpasswd MyS3cret -sqldblist Sales,Ledger
```

Babelfish Compass v.2023-06, June 2023

Compatibility assessment tool for Babelfish for PostgreSQL

Copyright Amazon.com, Inc. or its affiliates. All Rights Reserved.

Reading BabelfishFeatures.cfg

Latest Babelfish version supported: 3.2.0: BabelfishFeatures.cfg

Reading

C:\Users\johnsmith\Documents\BabelfishCompass\BabelfishCompassUser.Optimistic.cfg

Running Powershell/SMO script to generate DDL for server 'mybigbox' into directory

C:\Users\johnsmith\AppData\Local\Temp\CompassAutoDDL-2023-Jun-07-13.22.51...

Note: run times strongly depend on network proximity to the SQL Server.

Alternatively, generate DDL manually through SQL Server Management Studio on the SQL Server host.

You can abort by hitting CTRL-C

DDL generated in C:\Users\johnsmith\AppData\Local\Temp\CompassAutoDDL-2023-Jun-07-13.22.51 :

Volume in drive C is OSDisk

Volume Serial Number is FC25-F714

Directory of C:\Users\johnsmith\AppData\Local\Temp\CompassAutoDDL-2023-Jun-07-13.22.51

```
07-Jun-2023  13:24    <DIR>          .
07-Jun-2023  13:24    <DIR>          ..
07-Jun-2023  13:24                1,959,008 Ledger_SMO_DDL_2023-Jun-07.sql
07-Jun-2023  13:23                925,052 Sales_SMO_DDL_2023-Jun-07.sql
                2 File(s)            2,884,060 bytes
                2 Dir(s)  204,604,514,304 bytes free
```

Creating C:\Users\johnsmith\Documents\BabelfishCompass\SMOTest

Run starting : 07-Jun-2023 13:22:51 (Windows)

BabelfishFeatures.cfg file : v.3.2.0, Jun-2023

Target Babelfish version : v.3.2.0 (PG 15.3)

Command line arguments : SMOTest -sqlendpoint mybigbox -sqllogin sa -sqlpasswd

***** -sqldblist Sales,Ledger

[...rest of Compass session...]

Security

The Babelfish Compass tool is a stand-alone, on-premises program which does not store any confidential or sensitive information: all information stored is derived from the SQL/DDL scripts which the user provides as input.

The Babelfish Compass tool operates offline and does not perform any network access, with these exceptions:

- The **-pgimport** option, which connects to a PG instance.
- The automatic DDL generation with **-sqlendpoint**, **-sqllogin**, **-sqlpasswd**, which connects to a SQL Server instance.
- The automatic check-for-updates, which connects to GitHub (see below).

Other than in these cases, the Babelfish Compass tool makes no network connections that are invisible to the user. Also, Babelfish Compass does not "phone home".

The -pgimport option

The **-pgimport** option connects to a PostgreSQL instance and loads captured items into a database table.

Technically, Babelfish Compass creates files **pg_import.bat** (on Mac/Linux, **pg_import.sh**) and **pg_import.psql** in the **captured** directory. The **pg_import.bat/pg_import.sh** file executes **pg_import.psql**, which runs a **CREATE TABLE** and a **COPY** statement in PostgreSQL.

Babelfish Compass executes this function by spawning a subprocess to run **pg_import.bat/pg_import.sh**.

To make a connection to the PostgreSQL instance, the user must specify connection attributes on the Babelfish Compass command line, including the PostgreSQL username and password. These connection attributes are not written to any file, but are supplied as environment variables in the short-lived spawned subprocess. These environment variables are not accessible from outside the spawned subprocess.

Note that the connection attributes may be accessible through the command-line history in the command-line session that runs Babelfish Compass.

As for the uploaded captured items, it is assumed that the user owns the PostgreSQL instance and is responsible for granting access to the uploaded data.

The -sqlpasswd option

When using the **-sqlpasswd** option (see [Automatic DDL generation](#)), the specified password is passed on to the Powershell script. This script connects to the SQL Server specified with **-sqlendpoint**. The password itself is not stored or logged anywhere by Babelfish Compass. The **-sqlendpoint**

and **-sqllogin** options are included in the Compass report header under "Command line arguments". Note that the password and other connection attributes may be accessible through the command-line history in the command-line session that runs Babelfish Compass.

Automatic update check

Starting with Babelfish Compass version 2022-04, Compass checks whether a more recent version of itself is available at GitHub by making a REST call to GitHub repository. If so, it will print a message to inform the user, but not take any further action: the user must still download and install the update manually.

No information about these REST calls is collected or stored, other than by GitHub's default behavior.

To suppress the update check, specify the **-nouupdatechk** option.

Using Babelfish Compass to migrate to PostgreSQL

Babelfish Compass analyzes the SQL/DDL code for a SQL Server-based application for compatibility with Babelfish. The purpose of this analysis is to inform a Go/No Go decision about starting a migration project from SQL Server to Babelfish. For this purpose, Babelfish Compass produces an assessment report which lists (in great detail) all SQL features found in the SQL/DDL code, and whether or not these are supported by the latest version of Babelfish.

On a high level, the sequence of steps involved in a migration is as follows:

1. The application owner identifies the SQL Server databases required for the application that is considered for migration to Babelfish. The application owner must ensure there are no legal restrictions with respect to migrating the application in question.
2. Reverse-engineer the SQL Server database(s) in question with SQL Server Management Studio (SSMS). This is done in the SSMS Object Explorer by right-clicking a database and selecting **Tasks → Generate Scripts**, and following the dialog (making sure to turn on triggers, collations, logins, owners and permissions (turned off in SSMS by default), by clicking the **Advanced** button and turning on the respective options).
 - Babelfish Compass requires input scripts to be syntactically valid T-SQL, using **go** as a batch delimiter (i.e. **sqlcmd**-style scripts). Some tools may be able to reverse-engineer, but don't do this correctly or completely, or don't generate the required batch delimiters (like DBeaver). Therefore, we recommend using SSMS to generate a DDL script of the database(s).
3. SSMS produces a DDL/SQL script as output. Use this script (or scripts) as input for Babelfish Compass to generate an assessment report (see instructions and examples earlier in this User Guide).
4. Optionally, generate additional cross-reference reports to obtain additional details about the unsupported features.
5. Discuss the results of the Babelfish Compass assessment and interpret the findings in the context of the application to be migrated. In these discussions, it may be possible to descope the migration by identifying outdated or redundant parts of the application which do not need to be migrated.
6. Use the assessment results that show the unsupported SQL features in the SQL/DDL code, to decide if it is opportune to start a migration project to Babelfish. If the current version of Babelfish is deemed to be insufficiently compatible with the application in question, we

recommended you re-run the analysis when future releases of Babelfish are available which will provide more functionality.

7. If proceeding with a migration, modify the SQL/DDDL scripts to rewrite or remove the SQL/DDDL statements that are reported as **not supported** or **requiring review**. Then, invoke the SQL/DDDL script against Babelfish (with **sqlcmd**) to recreate the schema in Babelfish.
8. Finally, perform a data migration, and reconfigure the client applications to connect to Babelfish.

Please keep the following in mind:

- Admittedly, the amount of detail in a Babelfish Compass assessment report can be large. When discussing the Babelfish Compass findings with an application owner, make sure to highlight the many aspects that are supported by Babelfish: experience has shown that when focusing primarily on the non-supported features, SQL Server users may easily end up with an unnecessary negative perception of Babelfish's capabilities.
- A Babelfish migration involves more than just the server-side SQL/DDDL code, for example, interfaces with other system; ETL/ELT; SSIS/SSRS, replication, etc. These aspects may not be reflected in the server-side view provided by Babelfish Compass.

Troubleshooting

This section contains some troubleshooting tips. If you encounter unexpected behavior by Babelfish Compass, we recommend you first read this User Guide in detail.

- Syntax errors: while the SQL/DDDL input scripts are reverse-engineered from existing applications and their contents are assumed to contain syntactically valid SQL code, it is possible that the SQL code contains syntax errors. A syntax error might be the result of manual editing or inconsistencies. For example, MERGE statements must be terminated with a semicolon, while such a terminator is optional for most other T-SQL statements.

In case of a syntax error, this will be printed to **stdout**, and the offending batch will be logged in a file in the **errorbatches** subdirectory of the report directory. A batch containing syntax errors is not analyzed by Babelfish Compass.

Remedy: correct any SQL syntax errors and re-run the script through the Babelfish Compass tool using the **-replace** flag.

- If syntax errors are printed that show garbage characters, it may be that the input file encoding is not correctly specified. The default encoding, and all available encodings, are displayed with the **-encoding help** option.

Remedy: specify the correct encoding with **-encoding** on the command line.

Licensing

Copyright [Amazon.com](https://www.amazon.com), Inc. or its affiliates. All Rights Reserved.

SPDX-License-Identifier: Apache-2.0

GitHub: https://github.com/babelfish-for-postgresql/babelfish_compass